

W3 Experiments Summary: MLOps Pipeline for Medical Model Deployment

ChestMNIST — Full Fine-Tuning, Ablations & MLOps Deployment

Team F6-Score

Advanced Machine Learning Course, Spring 2026

National Higher School of Computer Science, Algiers, Algeria

Abstract—This report documents Week 3 deliverables for the MLOps Pipeline for Medical Model Deployment project. We implement the main method — full fine-tuning of MobileNetV2 and ResNet18 on ChestMNIST — and compare against Week 2 linear probe baselines and the Yang et al. [1] SOTA benchmark (macro-AUC 0.7707). MobileNetV2 achieves a test macro-AUC of 0.7900 (+0.105 over its linear probe baseline), while ResNet18 reaches 0.7820 (+0.096). Both models exceed SOTA. Two ablation experiments are conducted: per-class threshold tuning, which improves macro-F1 by +24% relative for MobileNetV2 (0.1884 $\rightarrow$ 0.2360), and an inference latency benchmark showing MobileNetV2 is $1.5\times$ faster than ResNet18 at 58.6 ms/image on CPU with $5\times$ fewer parameters. All experiments are tracked in MLflow and the best model is registered in the MLflow Model Registry, deployed via a FastAPI endpoint with A/B routing, and monitored via a live Streamlit dashboard implementing the Page-Hinkley drift detection test across four simulated hospital sites.

TABLE I
FINE-TUNING HYPERPARAMETERS

Parameter	Value
Loss	BCEWithLogitsLoss + pos_weight
Optimizer	Adam (differential LR)
LR backbone	1×10^{-4}
LR head	1×10^{-3}
Scheduler	ReduceLROnPlateau (factor=0.5, patience=2)
Batch size	64
Epochs	10
Input size	$224 \times 224 \times 3$
Seed	42
Train / Val / Test	78,468 / 11,219 / 22,433
Checkpoint criterion	Best validation macro-AUC

All runs are logged to MLflow under experiment `mlops-chestmnist` with per-epoch train loss, validation loss, macro-AUC, macro-F1, per-class AUC, and the best checkpoint stored as an MLflow artifact.

I. METHOD: FULL FINE-TUNING

A. Architecture

Both backbones follow an identical adaptation strategy. A `GrayscaleTo3Ch` module replicates the single-channel ChestMNIST input to 3 channels, preserving compatibility with ImageNet-pretrained weights without discarding the first convolutional layer. The classification head is replaced with a `Linear( $d$ , 14)` layer with sigmoid outputs for multi-label prediction: $d = 1280$ for MobileNetV2 and $d = 512$ for ResNet18. All backbone layers are fully unfrozen during training, in contrast to the frozen linear probe of Week 2.

B. Training Protocol

Differential learning rates are applied to prevent the randomly initialised head from distorting pretrained backbone features: the backbone receives $lr = 10^{-4}$ and the head receives $lr = 10^{-3}$. A `ReduceLROnPlateau` scheduler (factor=0.5, patience=2) monitors validation loss. Class imbalance is handled via `BCEWithLogitsLoss` with per-class `pos_weight` computed from the training set; without this, the model collapses to predicting all-negative for every class (F1 = 0.000). All other settings are identical to Week 2. Table I summarises the complete hyperparameter set.

II. EXPERIMENTAL PROTOCOL

The official MedMNIST train/val/test splits are preserved across all experiments (78,468 / 11,219 / 22,433 images). Seed 42 is fixed for all random operations (PyTorch, NumPy, Python). The primary metric is macro-AUC, chosen for its robustness to class imbalance and threshold independence. Macro-F1 at threshold 0.5 is reported as a secondary metric. Best checkpoints are selected by validation macro-AUC and evaluated *once* on the held-out test set to prevent test-set overfitting.

The learning curves (Fig. 1) show consistent improvement across 10 epochs with no early saturation — in contrast to Week 2 linear probes, which saturated at epoch 1. MobileNetV2 achieves its best validation AUC of 0.7850 at epoch 8, with test AUC 0.7900. ResNet18 peaks at epoch 7 (val AUC 0.7909), with test AUC 0.7820. The slight gap between best validation and test AUC is within normal variance.




Fig. 1. Validation macro-AUC per epoch for MobileNetV2 and ResNet18 during full fine-tuning (10 epochs). Both models improve steadily; neither saturates, confirming that backbone adaptation is necessary.

III. RESULTS: BASELINE VS. FINE-TUNED

A. Main Comparison

Table II reports all experimental conditions. Full fine-tuning yields +0.105 macro-AUC for MobileNetV2 and +0.096 for ResNet18 over their respective linear probe baselines. Both fine-tuned models exceed the Yang et al. SOTA of 0.7707 by +0.019 and +0.011 respectively.

TABLE II
FULL RESULTS SUMMARY (TEST SET)

Model	Strategy	AUC	F1@0.5	F1 tuned	Δ AUC
MobileNetV2	Linear probe	0.6852	0.1333	—	—
ResNet18	Linear probe	0.6865	0.1378	—	—
MobileNetV2	Full fine-tune	0.7900	0.1884	0.2360	+0.105
ResNet18	Full fine-tune	0.7820	0.1712	0.2331	+0.096
Yang et al. [1]	Full fine-tune	0.7707	—	—	—

B. Per-class AUC Analysis

Table III provides the complete per-class AUC breakdown across all four conditions. Fig. 2 visualises the progression from linear probe to fine-tuned for both backbones.

TABLE III
PER-CLASS AUC — ALL CONDITIONS (TEST SET)

Class	MV2-W2	MV2-W3	RN18-W2	RN18-W3
Atelectasis	0.8143	0.7726	0.8192	0.7745
Cardiomegaly	0.7566	0.8919	0.7521	0.8930
Effusion	0.7445	0.8445	0.7320	0.8510
Infiltration	0.7407	0.6809	0.7301	0.6829
Mass	0.7254	0.7739	0.7132	0.7545
Nodule	0.7106	0.6668	0.7052	0.6550
Pneumonia	0.6938	0.7228	0.6940	0.7163
Pneumothorax	0.6810	0.8281	0.6731	0.8210
Consolidation	0.6723	0.7844	0.6716	0.7823
Edema	0.6469	0.8718	0.6367	0.8672
Emphysema	0.6346	0.8255	0.6342	0.8178
Fibrosis	0.6118	0.7726	0.6315	0.7658
Pleural_Thickening	0.5814	0.7415	0.6228	0.7306
Hernia	0.5787	0.9113	0.5947	0.9272
Macro	0.6852	0.7900	0.6865	0.7820

Fine-tuning improves 12 of 14 classes for both backbones. The two exceptions — Atelectasis (−0.042 for MV2) and Infiltration (−0.060 for MV2) — are the two most prevalent classes in ChestMNIST, suggesting that the frozen linear probe had already learned a near-optimal linear separator for these high-frequency patterns, and that backbone adaptation slightly disrupts this.

The largest gains concentrate in rare, morphologically distinctive classes: Hernia (+0.333 for MV2), Edema (+0.225), Emphysema (+0.191), and Cardiomegaly (+0.135). These classes require deep spatial feature adaptation that a frozen ImageNet backbone cannot provide. Nodule (0.667) and Infiltration (0.681) remain the weakest classes for MV2, consistent with the 28×28 native resolution bottleneck: fine-grained texture differences are irretrievably lost during downsampling.




Fig. 2. Per-class AUC: linear probes (W2) vs. fine-tuned (W3) for both backbones. Fine-tuning improves 12 of 14 classes. Hernia shows the largest absolute gain (+0.333 for MV2).

C. SOTA Comparison

Fig. 3 plots our fine-tuned models against the Yang et al. benchmark. MobileNetV2 outperforms SOTA on 9 of 14 classes. Notably, SOTA uses ResNet18 full fine-tuning; our result shows MobileNetV2 achieves higher macro-AUC (0.7900 vs. 0.7707) with $5 \times$ fewer parameters (2.2M vs. 11.2M).



Fig. 3. Fine-tuned models vs. Yang et al. SOTA per-class AUC. MobileNetV2 (macro=0.7900) exceeds SOTA (0.7707) on 9 of 14 classes.

D. ROC Curves

Fig. 4 shows per-class ROC curves for the best model (MobileNetV2 fine-tuned). Hernia (AUC=0.911), Cardiomegaly (AUC=0.892), and Edema (AUC=0.872) exhibit near-ideal curves. Infiltration (0.681) and Nodule (0.667) show curves close to the diagonal, confirming the resolution bottleneck as a hard constraint on discriminability.




Fig. 4. Per-class ROC curves — MobileNetV2 fine-tuned (test set). Hernia, Cardiomegaly, and Edema are strongest; Infiltration and Nodule remain most challenging.

IV. ABLATION EXPERIMENTS

A. Ablation 1: Per-class Threshold Tuning

The default threshold of 0.5 is miscalibrated for severely imbalanced multi-label data: a model that never predicts positive achieves $F1 = 0$ but binary cross-entropy well below 0.5. We tune per-class thresholds on the *validation set* by sweeping $t \in [0.1, 0.9]$ at steps of 0.05 for each of the 14 classes independently, selecting the threshold maximising per-class F1. Tuned thresholds are applied at test time without retraining.

Table IV shows macro-F1 improves from 0.1884 to 0.2360 for MobileNetV2 (+0.046, +24% relative) and from 0.1712 to 0.2331 for ResNet18 (+0.062, +36% relative). The tuned thresholds in Table V are uniformly high (0.55–0.90), reflecting the aggressive positive-confidence required to overcome the prior towards all-negative predictions. AUC is unaffected by threshold choice, confirming it as the appropriate primary metric.

TABLE IV
ABLATION 1 — THRESHOLD TUNING (TEST SET MACRO-F1)

Model	F1 @ 0.5	F1 tuned	$\Delta F1$
MobileNetV2	0.1884	0.2360	+0.046 (+24%)
ResNet18	0.1712	0.2331	+0.062 (+36%)

TABLE V
TUNED THRESHOLDS — MOBILENETV2 (VAL SET OPTIMISED)

Class	Threshold
Atelectasis	0.70
Cardiomegaly	0.90
Effusion	0.65
Infiltration	0.55
Mass	0.80
Nodule	0.60
Pneumonia	0.85
Pneumothorax	0.80
Consolidation	0.75
Edema	0.90
Emphysema	0.80
Fibrosis	0.75
Pleural_Thickening	0.75
Hernia	0.90




Fig. 5. Ablation 1: per-class F1 at fixed threshold 0.5 vs. tuned threshold. Tuned thresholds improve every class.

B. Ablation 2: Inference Latency Benchmark

We benchmark both fine-tuned models on CPU (200 runs, 10 warmup iterations discarded) for single-image and batch-32 inference on 224×224 inputs. Table VI reports results.

TABLE VI
ABLATION 2 — INFERENCE LATENCY BENCHMARK (CPU, 200 RUNS)

Model	Params	ms/img	ms/b32	AUC
MobileNetV2	2.2M	58.6	1439.6	0.7900
ResNet18	11.2M	87.3	1820.4	0.7820

MobileNetV2 is $1.49\times$ faster per image (58.6 vs. 87.3 ms) and $1.26\times$ faster per batch, while achieving higher macro-AUC. With $5\times$ fewer parameters, MobileNetV2 is the unambiguous deployment choice: it dominates ResNet18 on accuracy, latency, and memory footprint simultaneously. At 58.6 ms per image on CPU, MobileNetV2 supports approximately 17 inferences per second on a single core, sufficient for real-time triage in a clinical radiology workflow.




Fig. 6. Ablation 2: inference latency comparison (left) and AUC vs. latency efficiency plot (right). MobileNetV2 dominates on both axes.

V. MLOPS PIPELINE

A. Experiment Tracking

All W2 and W3 runs are tracked in a single MLflow experiment (mlflow-chestmnist) stored in experiments/mlflow.db. Each run logs: backbone, strategy, all hyperparameters, per-epoch train/val loss, macro-AUC, macro-F1, per-class AUC, and the best checkpoint as an MLflow artifact. The database is committed to the repository; .pth checkpoints are excluded from Git and recoverable via scripts/restore_checkpoints.py, implementing the artifact-based reproducibility pattern of Zaharia et al. [5].




Fig. 7. MLflow experiment view (mlflow ui) showing all W2 and W3 runs with tracked metrics, parameters, and artifact links.

B. Model Registry and Promotion Gate

The best MobileNetV2 checkpoint is registered in the MLflow Model Registry via scripts/register_model.py, which enforces a two-gate validation policy before promotion (slide 21 of the course): the candidate must achieve macro-AUC ≥ 0.77 (Yang et al. SOTA threshold) and CPU single-image latency ≤ 200 ms. Both gates pass (AUC=0.7900, latency=58.6 ms), and the model is assigned the production alias (Fig. 8), implementing the governed artifact lifecycle: None $\rightarrow$ Staging $\rightarrow$ Production [5]. This separation between experimentation and deployment is a core MLOps maturity requirement [4].




Fig. 8. MLflow Model Registry showing ChestMNIST-MobileNetV2 v1 with the production alias, assigned after passing the AUC and latency promotion gates.

C. Model Serving

The production checkpoint (MobileNetV2, AUC=0.7900) is served via a FastAPI application (src/api.py) exposing three endpoints: GET /health returns model load status; POST /predict/{model_name} accepts an uploaded image and returns per-class probabilities, positive predictions, applied thresholds, and inference latency; POST /predict/ab implements randomised A/B routing between MobileNetV2 and ResNet18, simulating canary deployment for live model comparison. The tuned thresholds from Ablation 1 are applied at the MobileNetV2 endpoint. The preprocessing pipeline is identical to the training transform (Resize $\rightarrow$ Grayscale $\rightarrow$ ToTensor $\rightarrow$ Normalize), preventing training-serving skew [4]. All inference requests are appended to logs/inference_log.jsonl with timestamp, model name, latency, and prediction metadata, providing a full audit trail for post-deployment analysis.




Fig. 9. FastAPI Swagger UI (/docs) showing the POST /predict/mobilenet endpoint. The response includes per-class probabilities, positive labels, threshold mode, and latency in milliseconds.

D. Reproducibility

The complete pipeline is fully reproducible from the repository: (1) pip install -r requirements.txt, (2) python scripts/restore_checkpoints.py recovers all .pth checkpoints, (3) python scripts/export_test_images.py generates hospital simulation images, (4) python

scripts/register_model.py registers and promotes the model, (5) uvicorn src.api:app --port 8000 starts the serving endpoint, and (6) http://localhost:8000/docs#/ to test the endpoints.

VI. CRITICAL ANALYSIS

A. Interpretation of Results

Full fine-tuning is necessary. The +0.105 AUC gain confirms that ImageNet features require full adaptation for medical imaging. The Week 2 saturation at epoch 1 is completely resolved: training loss decreases monotonically over 10 epochs for both models, indicating stable gradient flow through the unfrozen backbone.

MobileNetV2 is strictly preferable. It achieves higher macro-AUC (0.7900 vs. 0.7820), lower latency (58.6 vs. 87.3 ms/img), and $5\times$ fewer parameters. There is no accuracy justification for deploying ResNet18 in this setting.

Macro-F1 is recoverable without retraining. The low F1 at threshold 0.5 (0.1884) is a calibration artefact of class imbalance, not a model failure. Per-class threshold tuning on the validation set recovers a 24% relative improvement (to 0.2360) at zero additional training cost. Deployment thresholds should always be tuned on local validation data, not inherited from training defaults.

Both models exceed SOTA. Our MobileNetV2 (0.7900) and ResNet18 (0.7820) surpass Yang et al. [1] (0.7707) by +0.019 and +0.011 respectively. The SOTA reference uses ResNet18 full fine-tuning; our result demonstrates that a lighter architecture achieves higher accuracy with lower operational cost.

B. Limitations

Resolution bottleneck. Native ChestMNIST resolution is 28×28 . Upscaling to 224×224 does not recover lost high-frequency information. Nodule (AUC=0.667) and Infiltration (AUC=0.681) are fundamentally limited by this.

CPU-only latency. GPU latency (expected <5 ms/image) remains for future work.

No uncertainty quantification. The model produces point estimates. Clinical deployment would require uncertainty-aware predictions for radiologist handoff.

REPRODUCIBILITY SUMMARY

Training: notebooks/W3_finetuning.ipynb

Evaluation: notebooks/W3_evaluation.ipynb

MLflow UI: mlflow ui --backend-store-uri sqlite:///experiments/mlflow.db

Registry: python scripts/register_model.py

API: uvicorn src.api:app --port 8000

Checkpoints: python scripts/restore_checkpoints.py

Test images: python scripts/export_test_images.py

REFERENCES

- [1] J. Yang et al., “MedMNIST v2: A Large-Scale Lightweight Benchmark for 2D and 3D Biomedical Image Classification,” *Scientific Data*, vol. 10, no. 1, 2023.
- [2] M. Sandler, A. Howard, M. Zhu, A. Zhmoginov, and L.-C. Chen, “MobileNetV2: Inverted Residuals and Linear Bottlenecks,” *CVPR*, 2018.
- [3] K. He, X. Zhang, S. Ren, and J. Sun, “Deep Residual Learning for Image Recognition,” *CVPR*, 2016.
- [4] D. Sculley et al., “Hidden Technical Debt in Machine Learning Systems,” *NeurIPS*, vol. 28, 2015.
- [5] M. Zaharia et al., “Accelerating the Machine Learning Lifecycle with MLflow,” *IEEE Data Engineering Bulletin*, vol. 41, no. 4, pp. 39–45, 2018.
- [6] A. Bifet and R. Gavalda, “Learning from Time-Changing Data with Adaptive Windowing,” *SDM*, 2007.
- [7] J. Gama, P. Medas, G. Castillo, and P. Rodrigues, “Learning with Drift Detection,” *SBIA*, 2004.